



POLITECNICO DI TORINO

Master degree in Cybersecurity

Master Degree Thesis

An incremental approach to automatic firewall configuration

Supervisor

prof. Fulvio Valenza
prof. Daniele Bringhenti
prof. Francesco Pizzato

Candidate

Filippo DEL MINISTRO

ACADEMIC YEAR 2025-2026

Nowadays, computer networks are complex systems. Modern networks are in fact often distributed systems in which security policy enforcement is a critical and continuous operational challenge. As topologies grow in scale and diversity, the manual configuration of packet-filtering firewalls becomes increasingly prone to errors. Misconfigurations can expose sensitive resources or degrade service availability. In this context, the shift toward microsegmentation, distributing enforcement across multiple firewall instances throughout the network, improves security granularity but amplifies the management effort. Automated tools that generate formally correct and optimal configurations are therefore essential.

VEREFOO (VERification REasoning on Firewall Optimization) addresses this need by automatically computing firewall allocation and configuration from a network topology and a set of Network Security Requirements (NSRs). It encodes the problem as a formal optimisation problem — specifically a Maximum Satisfiability Modulo Theories (MaxSMT) instance. This approach provides two key guarantees: formal correctness, meaning that the generated configuration provably satisfies every specified requirement, and minimality, meaning that the number of allocated firewall instances and configured rules is kept as low as possible. However, VEREFOO follows a monolithic approach: every time the NSR set changes, the entire problem is re-encoded and solved from scratch, scaling poorly as the network or requirement set grows.

React-VEREFOO addresses this by taking an incremental approach. It identifies which requirements are new, removed, or unchanged, and restricts the solver to only the affected parts of the network. This reduces reconfiguration cost when few requirements change, but the benefit decreases when a large fraction of the NSR set is modified simultaneously.

This thesis pushes the incremental principle to its extreme. The proposed Iterative React-VEREFOO introduces new NSRs one at a time: at each step, the current configuration is used as the starting state and exactly one requirement is added. React-VEREFOO computes the minimal update, and the result becomes the input to the next iteration. By construction, the reconsidered portion of the network is always minimal, keeping each MaxSMT subproblem small regardless of the total number of requirements.

The approach is evaluated on a set of benchmark topologies covering both real-world and synthetic networks of varying scale, with the aim of quantifying whether the consistent reduction in per-step problem size translates into a measurable improvement in total execution time compared to the monolithic VEREFOO baseline.

Index

1	Introduction	7
2	Background	9
2.1	Context	9
2.2	VEREFOO	10
2.2.1	Overview	10
2.2.2	Network Security Requirements	12
2.2.3	Service Graph	13
2.2.4	VEREFOO Output	14
2.2.5	VEREFOO Validation	14
2.3	React-VEREFOO	14
3	Thesis Objectives	18
3.1	Idea of the Solution	18
3.2	Expected Benefits and Trade-offs	20
4	The Iterative VEREFOO Approach	22
4.1	Algorithm Description	22
4.1.1	Vanilla VEREFOO	22
4.1.2	React-VEREFOO	23
4.1.3	Iterative VEREFOO	24
4.1.4	Structural Comparison	24
4.2	Worked Example	25
4.2.1	Network Topology and NSR Set	25
4.2.2	Iteration 1 — Introducing p_1	26
4.2.3	Iteration 2 — Introducing p_2	26
4.2.4	Iteration 3 — Introducing p_3	27
4.2.5	Final Configuration and Summary	27

5	Implementation and Validation	29
5.1	Implementation	29
5.1.1	Entry Point and initialization	29
5.1.2	Core change: the Iterative Loop	30
5.1.3	Summary of Changes	31
5.2	Testing	31
5.2.1	Qualitative Validation on CESNET	31
5.2.2	A Closer Look: Tie-Breaking and Flow Analysis	40
5.2.3	Discussion and Limitations	40
6	Conclusion and Future Work	44
6.1	Conclusions	44
6.2	Future Work	44
	Bibliography	46

Chapter 1

Introduction

Modern enterprise networks are large, dynamic infrastructures that change continuously. New services are deployed, nodes are added or removed, and connectivity requirements evolve at a pace that security policy management must match. At the same time, these networks reach hundreds of heterogeneous nodes distributed across sites and organisational boundaries, making security enforcement a complex and ongoing operational challenge [1]. Packet-filtering firewalls are the primary mechanism for enforcing security policies, and their correct configuration is essential to prevent unauthorized access, data breaches, and service disruptions.

The traditional approach to firewall configuration relies on manual translation of high-level security requirements into device-level rules, and this approach does not scale to modern network environments. As networks grow in size and change more frequently, the volume and complexity of rules to be managed quickly exceeds human capacity. More critically, manual processes offer no formal guarantee of correctness: a single misconfiguration may silently violate security requirements or block legitimate traffic, with consequences that can remain undetected for extended periods.

VEREFOO was developed to address this problem through full automation, ensuring formal correctness and optimality. Given a network topology and a set of Network Security Requirements (NSRs), it automatically determines where firewalls should be placed and what filtering rules each should enforce [2]. It achieves this by encoding the problem as a formal optimisation problem, namely a Maximum Satisfiability Modulo Theories (MaxSMT) instance, solved by the Z3 solver. However, VEREFOO follows a monolithic strategy: any change to the NSR set triggers a complete recomputation from scratch, making it poorly suited to dynamic environments where requirements change frequently such as modern enterprise networks.

React-VEREFOO addressed this limitation by adopting an incremental approach [3]. Rather than discarding the existing configuration, it identifies which NSRs are new, which have been removed, and which are unchanged, and restricts the solver to only the affected parts of the network. This yields significant speedups when the number of modified requirements is small. However, when a large fraction of the NSR set changes simultaneously, the advantage in terms of execution time is reduced and performance results similar to the monolithic baseline.

This thesis proposes a strategy that preserves the benefits of the incremental approach regardless of how many novel requirements are introduced. The key idea is to push the incremental principle to its extreme: the iterative approach proposes to perform a full configuration from scratch by introducing new NSRs one at a time, invoking React-VEREFOO internally at each step to compute the minimal update on the current configuration. By construction, the modified fraction at each step is always minimal — one single requirement — so the subproblem remains small and the performance advantage of the incremental approach is consistently in effect.

The central research question this thesis aims to answer is whether invoking React-VEREFOO multiple times during the iterative approach yields a lower total execution time than vanilla VEREFOO, which solves the same problem in a single monolithic call. The intuition is that many cheap solver calls may collectively outperform a single expensive one; whether this holds in practice, and under what conditions, is the empirical question addressed in Chapter 5. The full formulation of the approach and the thesis objectives are presented in Chapters 3 and 4.

The remainder of the thesis is organised as follows. Chapter 2 provides the necessary background on network security automation, VEREFOO, and React-VEREFOO. Chapter 3 states the thesis objectives and introduces the iterative approach at a high level. Chapter 4 develops the approach from a theoretical perspective. Chapter 5 describes the implementation and reports the experimental validation of the approach. Chapter 6 draws conclusions and outlines directions for future work.

Chapter 2


Background

2.1 Context

Network security is a critical concern for almost every organization, from small businesses to large enterprises and public institutions. All of them rely on computer networks to support their operations, and all of them are exposed to the risk of unauthorized access, data breaches, and service disruptions. As a result, enforcing a well-defined and correct security policy is not optional, but instead is a fundamental operational requirement.

Modern enterprise networks are characterized by growing complexity: hundreds of interconnected nodes, heterogeneous services, and security policies that must be consistently enforced across distributed infrastructure. Moreover, network topologies are not static: new services are deployed, nodes are added or removed, and connectivity requirements evolve continuously, making security management an ongoing and demanding task. The cost of managing network security in this context is significant, as it requires specialized expertise, advanced planning, and constant monitoring.

A central challenge in this landscape stems from the evolution of network security architectures. Traditionally, a single perimeter firewall was sufficient to protect a network by controlling traffic at its boundary. However, as networks grew in scale and complexity, this model was proved inadequate: a single point of control cannot enforce fine-grained policies across a differentiated internal topology. Modern networks have therefore evolved toward microsegmentation, deploying multiple packet-filtering firewall instances strategically distributed across the topology to enforce security policies at a much more granular level, as illustrated in Figure 2.1. This distributed approach offers clear advantages such as enabling finer-grained traffic control, reducing single points of failure, and allowing security policies to be enforced closer to the traffic sources. On the other hand, this approach also introduces considerable management complexity. Each firewall instance must be individually configured with a set of filtering rules, and those rules must collectively enforce the intended security policy across the entire network, without conflicts or gaps.



(a) Traditional perimeter firewall architecture. A single point of control at the network boundary; internal traffic is unfiltered.

(b) Distributed firewall architecture (microsegmentation). A filtering instance guards each segment, so internal traffic is filtered as well.

Figure 2.1: Evolution of network security architectures: from a single perimeter firewall to distributed packet-filtering instances deployed across the topology.

In this context, the configuration of firewalls becomes a critical task. On one hand, overly permissive rules may allow unauthorized traffic to reach sensitive resources, exposing the network to attacks. On the other hand, overly restrictive rules may block legitimate traffic, degrading service availability and causing operational disruptions. Both failure modes carry significant consequences, and yet both are commonly caused by a manual configuration process that is inherently prone to errors, especially for large-scale networks [4].

The gap between the complexity of modern networks and the limitations of manual security management calls for automated approaches that can generate correct and optimal firewall configurations without human intervention. It is in this context that VEREFOO was developed: a tool for the automated, formally correct allocation and configuration of distributed firewalls, designed to bridge exactly this gap.

2.2 VEREFOO

2.2.1 Overview

VEREFOO (VERification REasoning on Firewall Optimization) is a tool designed to automate the allocation and configuration of packet-filtering firewalls in a network, eliminating entirely the need for manual intervention [2]. As exemplified in Figure 2.2, given a network topology and a set of Network Security Requirements (NSRs) — high-level specifications describing which traffic flows must be allowed or denied between network endpoints [5] — VEREFOO automatically determines where firewalls should be placed within the topology and generates the corresponding filtering rules for each allocated instance. The resulting configuration is expressed in a vendor-independent format, which can then be translated into concrete rules for specific implementations such as iptables, BPF, or OVS.

The tool provides three formal guarantees. *Automation* ensures that no human intervention is required at any stage of the process. *Formal correctness* guarantees that the generated configuration satisfies all specified NSRs by construction, eliminating the risk of undetected misconfigurations. *Optimality* ensures that both the number of allocated firewall instances and the number of generated rules are minimised, reducing resource consumption and improving performance. These properties are achieved by encoding the firewall allocation and configuration problem as a partial weighted Maximum Satisfiability Modulo Theories (MaxSMT) instance, which is then solved using the Z3 SMT solver. The hard constraints of the problem encode the correctness requirements, while the soft constraints encode the optimality objectives; the solver finds a solution that satisfies all hard constraints and maximises the sum of satisfied soft constraint weights.




Figure 2.2: High-level view of VEREFOO: given the Service Graph and the Network Security Requirements, the tool produces the firewall allocation scheme and the filtering policy for each allocated instance guaranteeing formal correctness and optimality.

Before the MaxSMT problem is formulated, VEREFOO must compute the set of traffic flows to be modelled. Two algorithms are available for this step. The *Maximal Flows* algorithm (MF) or *Atomic Predicates* algorithm (AP). The MF algorithm groups packets into maximal flows, which are flows that are not sub-flows of any other flow; in other words, multiple packets with identical traversal behaviour are represented by a single entity, reducing the number of constraints in the problem [2]. The Atomic Predicates algorithm instead goes further, decomposing every traffic predicate into a set of atomic predicates that are minimal, unique, and mutually disjoint, then expressing each flow in terms of these atomic predicates. This additional decomposition reduces the number of distinct cases the solver must consider and has been shown to provide better performance for automatic (re)configuration problems [3].

2.2.2 Network Security Requirements

Network Security Requirements (NSRs) are the high-level specification of the security policy that a network must enforce. Rather than describing firewall rules directly, NSRs express what the administrator intends: which traffic flows between pairs of endpoints must be permitted and which must be denied. This level of abstraction decouples the intent of the security policy from its low-level implementation, allowing VEREFOO to take full responsibility for translating requirements into concrete filtering rules.

Each NSR refers to a specific traffic flow, identified by five attributes: source address, destination address, transport protocol, source port, and destination port. The administrator specifies, for each such flow, whether it must be reachable (i.e., allowed end-to-end through the network) or isolated (i.e., blocked). VEREFOO supports four different policy modes that determine how the requirements are interpreted: *whitelisting*, where all traffic is blocked by default and only explicitly allowed flows are permitted; *blacklisting*, where all traffic is allowed by default and only explicitly listed flows are denied; and two mixed modes (*rule-oriented specific* and *security-oriented specific*), which accept both reachability and isolation requirements without a fixed default behavior, differing in how unspecified flows are handled [2].

NSRs are provided to VEREFOO as part of its input, alongside the network topology. Internally, they are represented and stored as a structured set of logical constraints, expressed in a vendor-independent intermediate language. A higher-level specification language is also supported, from which the corresponding medium-level NSRs are automatically derived through a translation step. VEREFOO processes this set together with the Allocation Graph derived from the topology, and encodes each NSR as a hard constraint in the MaxSMT problem: the solver is required to find a firewall allocation and rule set that satisfies all of them without exception. If no valid solution exists — for example because user-defined constraints on the Allocation Graph have excluded all possible firewall positions — VEREFOO produces a non-enforceability report rather than a configuration.

Table 2.1 shows three representative NSRs illustrating how requirements of different types are expressed.

Table 2.1: Representative NSR examples.

Type	Src	Dst	Protocol	Port	Description
Reachability	40.40.41.1	130.10.0.1	any	any	trusted client may reach server
Isolation	40.40.42.1	130.10.0.1	any	any	untrusted client is blocked
Reachability	40.40.41.1	130.10.0.1	TCP	80	allow TCP traffic on port 80

2.2.3 Service Graph

A Service Graph (SG) is the logical representation of a virtual network: a directed graph in which nodes represent Network Functions (NFs) such as web servers, load balancers, caches, or client subnetworks while edges represent the connectivity between them. The SG captures all possible end-to-end traffic paths through the network at an abstract level, without exposing low-level forwarding infrastructure such as switches and routers. It is defined taking into consideration solely the services that the network must provide, independently of any security consideration [2].

The SG serves as the structural input to VEREFOO: it defines the space of possible firewall positions and the traffic paths that any allocated firewall must be able to intercept. To make this explicit, VEREFOO automatically derives an internal representation called the *Allocation Graph* (AG) from the SG. For each link between any pair of nodes in the SG, the AG introduces a placeholder element called an *Allocation Place* (AP) which represents a candidate position where a firewall instance may be allocated. Those Allocation Places are then used by VEREFOO during the optimal placement of firewalls. Optionally, VEREFOO allows manually constraining this process by forcing a firewall into a specific AP or excluding certain APs from consideration, for example to model pre-existing hardware firewalls that cannot be moved. Figure 2.3 illustrates how a simple Service Graph is transformed into its corresponding Allocation Graph.

VEREFOO uses the Allocation Graph (AG) jointly with the Network Security Requirements (NSRs) to formulate the MaxSMT problem. Each AP corresponds to a boolean decision variable that the solver uses to determine the optimal firewall placement and what filtering rules each firewall should enforce, such that all NSRs are satisfied and the total number of allocated firewalls and rules is minimised.




Figure 2.3: From Service Graph (left) to Allocation Graph (right). VEREFOO automatically derives the Allocation Graph by inserting an Allocation Place (shown as dashed boxes) on each link, representing a candidate position where a firewall instance may be allocated by the solver.

2.2.4 VEREFOO Output

When the MaxSMT problem is solved successfully, VEREFOO produces two complementary outputs [2]. The first is the *firewall allocation scheme*: an updated version of the Allocation Graph in which each selected Allocation Place hosts an active firewall instance, with the minimum number of firewalls necessary to enforce all NSRs. The second is the *Filtering Policy* (FP) for each allocated firewall: a complete, auto-generated rule set that defines how the firewall must handle every traffic flow passing through it. Each FP is expressed in a vendor-independent abstract language, making it independent of any specific firewall implementation. It consists of a default action (either drop-all in whitelisting mode, or allow-all in blacklisting mode) and a minimal set of explicit rules covering the traffic flows specified by the NSRs. Both the default action and the explicit rules are guaranteed to be anomaly-free, with no conflicts or redundancies among them.

Together, these two outputs constitute the full security configuration of the network at the logical level of the Service Graph, independent of the physical infrastructure. A subsequent translation step converts the abstract filtering policies into concrete rules for specific implementations such as iptables, BPF, or OVS [2]. If instead no solution exists, for example because user-defined constraints on the AG have excluded all viable firewall positions, VEREFOO returns a non-enforceability report in place of the configuration.

2.2.5 VEREFOO Validation

VEREFOO has been validated on both synthetic and real-world network topologies, including well-known benchmarks such as the Stanford and CESNET networks, demonstrating reliability and scalability across a range of network sizes and requirement sets [2].

Despite these strengths, VEREFOO’s monolithic approach has a fundamental scalability limitation: the entire set of NSRs and the full network topology are encoded into a single MaxSMT problem. As the number of requirements or the size of the network grows, the complexity of the problem increases significantly, leading to prohibitive computation times. This limitation motivates the development of more efficient alternatives, starting with React-VEREFOO.

Further details on the experimental validation of VEREFOO and its comparison with the iterative approach proposed in this thesis can be found in Chapter 5.

2.3 React-VEREFOO

In operational environments, network security policies are not static: NSRs evolve over time as new services are deployed, topology changes occur, or threat models are updated. A fundamental limitation of the monolithic approach used by VEREFOO is that any change to the NSR set triggers a full recomputation of the firewall configuration from scratch. React-VEREFOO addresses this by offering a more efficient approach in the dynamic context described above, avoiding the need to

discard a valid existing configuration when only a small portion of the requirements has changed.

React-VEREFOO takes two inputs: the existing firewall configuration (allocation scheme and filtering rules) and a pair of NSR sets — the *initial set* R_i , describing the requirements already enforced by the current configuration, and the *target set* R_t , describing the requirements that must be enforced after reconfiguration. The approach proceeds in three steps.




Figure 2.4: High-level view of React-VEREFOO: given the existing configuration, the initial and the updated NSR sets, it classifies requirements into added, kept, and deleted groups, identifies the minimal network area affected by the added requirements, and solves a reduced MaxSMT problem restricted to that area, leaving the rest of the configuration untouched.

Step 1 — NSR classification. React-VEREFOO partitions every NSR into one of three groups:

- $R_d = \{r \in R_i \mid r \notin R_t\}$: *deleted* requirements, previously enforced but no longer needed. They are simply excluded from the new problem formulation.
- $R_a = \{r \in R_t \mid r \notin R_i\}$: *added* requirements, new NSRs that must be enforced but are not yet satisfied. These are the sole driver of reconfiguration.
- $R_k = \{r \in R_t \mid r \in R_i\}$: *kept* requirements, already correctly enforced and still valid. The nodes satisfying them are left unchanged.

Step 2 — Identification of the network area to reconfigure. For each added requirement in R_a , React-VEREFOO analyses the Allocation Graph to determine which nodes are in conflict with the new requirement and therefore need to be reconsidered. The procedure differs depending on whether the added requirement is an isolation or a reachability constraint [3].

For an added *isolation* requirement (traffic that must be blocked), the algorithm examines all atomic flows associated with that requirement. If a flow is already blocked by an allocated firewall, no action is needed for that flow. If instead no node along the path currently blocks the traffic, all nodes on that path are flagged as eligible for reconfiguration, since one of them must be configured to enforce the isolation.

For an added *reachability* requirement (traffic that must be allowed), the algorithm checks whether at least one atomic flow reaches the destination without being blocked. If such a flow already exists, the requirement is already satisfied and no reconfiguration is needed. Otherwise, the nodes that are currently blocking the traffic are flagged as eligible for reconfiguration.

In both cases, the algorithm selects all nodes that *could potentially* be used to fulfill the new requirement, since the optimal choice among them is left to the solver. Nodes not flagged by any added requirement retain their current configuration and are excluded from the decision variables of the subsequent MaxSMT problem.

An example of this behaviour for both cases is shown in Figure 2.5. The topologies in this example are purely illustrative and should not be considered realistic network scenarios; they are presented solely to clarify what happens internally when a new isolation or reachability requirement is introduced.




Figure 2.5: React-VEREFOO Step 2: example of identification of the network area to reconfigure, depending on the type of added NSR. **(a) New isolation requirement:** both FW₁ and FW₂ are initially configured to allow Client→Server traffic. Since no node currently blocks the flow, all nodes on the path are flagged as candidates for reconfiguration, and one of them must be configured to enforce the new requirement. **(b) New reachability requirement:** FW₁ is initially configured to deny Client→Server traffic, while FW₂ allows it. Since FW₁ is the node actively blocking the flow, only it is flagged for reconfiguration while FW₂ is left unchanged.

Step 3 — MaxSMT problem formulation and solving. Once the reconfigurable area has been identified, React-VEREFOO formulates a MaxSMT problem. This problem models traffic flows for all NSRs in the target set, ensuring that the final configuration is correct with respect to all requirements. However, only the

flagged nodes appear as free decision variables; the configuration of all other nodes is treated as fixed. This restriction dramatically reduces the size of the search space compared to a full recomputation.

The soft constraints are also adjusted to bias the solver toward reusing the existing configuration: allocating an already-deployed firewall is preferred over placing a new one, and retaining an existing rule is preferred over generating a new one. This minimises the number of changes applied to the network, reducing the time needed to deploy the new configuration.

The result is an updated allocation scheme and a revised set of filtering policies that satisfy all requirements in the target set, with formal correctness guaranteed by construction. As with VEREFOO, the output may not be globally optimal, since only a subset of the network is reconsidered, but it is optimal within the identified reconfiguration area.

Experimental results [3] confirm that React-VEREFOO is significantly more efficient than vanilla VEREFOO when the proportion of modified NSRs is small, as the reduced problem size leads to much shorter solving times. However, when a large fraction of the requirements changes, the benefit is reduced and performance converges toward the monolithic baseline. This observation motivates the approach investigated in this thesis: by pushing the incremental logic to its extreme and introducing only one new NSR per iteration. The proportion of modified requirements is always minimised, potentially achieving consistent speedups regardless of the total number of requirements.

Chapter 3

Thesis Objectives

3.1 Idea of the Solution

To understand the motivation of this thesis, it is necessary to examine how React-VEREFOO determines what to reconfigure. Given an initial set of NSRs R_i already enforced in the network and a target set R_t that must be enforced after reconfiguration, React-VEREFOO classifies every requirement into one of three groups [3]:

- $R_d = \{r \in R_i \mid r \notin R_t\}$: *deleted* requirements, present in the initial configuration but no longer needed;
- $R_a = \{r \in R_t \mid r \notin R_i\}$: *added* requirements, new requirements that must be enforced but are not yet satisfied;
- $R_k = \{r \in R_t \mid r \in R_i\}$: *kept* requirements, already enforced and still valid in the target configuration.

Only the added requirements R_a drive the reconfiguration: React-VEREFOO identifies which nodes in the Allocation Graph are in conflict with each added NSR and restricts the solver to those areas, while leaving the rest of the configuration untouched. The final MaxSMT problem models all traffic flows for the target set $R_t = R_a \cup R_k$, but the decision variables are limited to the nodes identified as needing reconfiguration, making the problem significantly smaller than a full recomputation.

The key performance parameter is the proportion of kept requirements, which can be defined as:

$$\text{PercReqKept} = \frac{R_k}{R_k + R_a + R_d}$$

The higher this ratio, the smaller the added and deleted sets, and consequently the narrower the network area that must be reconsidered. Experimental results in [3] confirm that the approach yields its greatest speedups at high values of PercReqKept, and that performance converges toward the vanilla baseline as PercReqKept decreases. This is consistent with real-world operational patterns: studies of large-scale infrastructure [3] report that the vast majority of configuration updates affect

only a small fraction of the total rule set, making high PercReqKept the typical case in practice.

The key idea of this thesis is to push this observation to its logical extreme. Instead of processing all new NSRs in a single reconfiguration step, the proposed approach introduces them **one at a time**. At each step, only one new requirement is presented to React-VEREFOO as added, while everything enforced so far falls into the kept set. This means the reconfiguration area identified at each step is always the smallest it could possibly be — a single requirement drives each solver call, regardless of how many requirements must ultimately be incorporated. The proportion of kept requirements is therefore always at its theoretical maximum for the given state of the network, and grows larger with every iteration as the kept set expands.

Concretely, the process unfolds as a chain of reconfiguration steps. The network starts from an existing valid configuration. One new requirement is then selected and processed by the iterative system, which internally invokes React-VEREFOO to compute the minimal update to the current configuration. The result, the new allocation scheme and updated filtering rules become the starting point for the next step. Another requirement is then introduced, the process repeats, and so on until every new requirement has been incorporated. At no point is the full set of new requirements processed at once; the tool always sees exactly one addition on top of an otherwise stable configuration.

Formally, let P_0 denote the initial set of requirements already enforced in the network and $\{A_1, A_2, A_3, \dots\}$ the sequence of new NSRs to be introduced. The iterative process unfolds as the following steps:

1. initial: P_0 , additional: $\{A_1\}$
2. initial: $\{P_0, A_1\}$, additional: $\{A_2\}$
3. initial: $\{P_0, A_1, A_2\}$, additional: $\{A_3\}$
4. ...

The total number of solver invocations equals the number of new NSRs — a linear cost — but each individual invocation is far cheaper than a single reconfiguration over the full batch, since the solver always operates considering one new requirement only. Whether the cumulative saving per step outweighs the cost of repeating the process n times is the central empirical question addressed in Chapter 5.




Figure 3.1: The iterative approach: each step introduces exactly one new NSR, so the solver always sees the smallest possible incremental problem.

3.2 Expected Benefits and Trade-offs

The proposed approach is expected to yield consistent performance improvements over both vanilla VEREFOO and standard React-VEREFOO, at the cost of some trade-offs that are important to acknowledge.

Expected benefits. Because each invocation of React-VEREFOO sees a problem with exactly one new requirement, the solver is always operating on the smallest possible incremental problem. This is in contrast to standard React-VEREFOO, which can receive an arbitrarily large batch of new NSRs and whose efficiency degrades as the batch grows. The proposed iterative approach removes this dependency on batch size: regardless of how many NSRs must be processed in total, each individual solver call remains minimal.

Trade-offs. The main cost of the approach is the linear number of solver invocations: where vanilla VEREFOO solves one problem and React-VEREFOO solves one (possibly large) problem, the iterative variant proposed in this thesis solves n small problems. Each invocation carries a fixed overhead due to parsing, constraint encoding, and solver initialization, which accumulates over the n iterations. Whether solving n small problems is globally faster than solving one large problem, is the central question this thesis aims to address.

A second trade-off concerns the quality of the final configuration. Vanilla VEREFOO optimises globally over the full NSR set, producing a configuration that minimises the total number of firewalls and rules jointly. The proposed iterative approach, by contrast, optimises locally at each step. In particular, a firewall placed during the first iteration to satisfy the first new NSR (defined as A_1 in previous example) may not be the best structural choice once the subsequent new requirements ($A_2, \dots, A_n$) are taken into account in the following iterations. Whether and to what extent this affects the quality of the final configuration in practice remains an open question, and could be the subject of further investigation.

The approach investigated in this thesis focuses exclusively on the addition of new NSRs. The case of NSR removal is not considered: when a requirement must

be revoked, the iterative strategy described here does not apply, and a separate reconfiguration procedure would be needed. Similarly, the thesis considers only the whitelisting and blacklisting policy modes supported by React-VEREFOO, and the evaluation is conducted on three specific benchmark topologies. Generalisation to other policy modes or network types is left as future work.

Chapter 4

The Iterative VEREFOO Approach

4.1 Algorithm Description

This section describes the iterative approach in detail, highlighting the structural differences with respect to VEREFOO vanilla. Both variants share the same underlying MaxSMT engine and the same formal correctness guarantees; what changes is how the set of NSRs is presented to the solver across invocations.

4.1.1 Vanilla VEREFOO

In the standard, non-incremental version of VEREFOO, all Network Security Requirements are collected into a single set and passed to the solver in one call. The algorithm proceeds as follows.

Algorithm 1 Vanilla VEREFOO

Require: Service Graph G , full NSR set $P = \{p_1, p_2, \dots, p_n\}$

Ensure: Firewall allocation and filtering rules satisfying all $p \in P$, or UNSAT

- 1: Build Allocation Graph from G
 - 2: $result \leftarrow \text{SOLVEMAXSMT}(G, P)$
 - 3: **if** $result = \text{SAT}$ **then**
 - 4: Extract and return firewall allocation and filtering rules
 - 5: **else**
 - 6: **return** UNSAT
 - 7: **end if**
-

A single MaxSMT call encodes all n requirements simultaneously. The solver searches for a globally optimal configuration, minimising the total number of allocated firewalls and rules across the entire NSR set.

At the implementation level, vanilla VEREFOO is invoked through the following constructor of `VerefooSerializer`, which takes the network topology and the algorithm variant as inputs:

Listing 4.1: Vanilla VEREFOO constructor.

```
public VerefooSerializer(NFV root, String algo) { ... }
```

4.1.2 React-VEREFOO

React-VEREFOO extends vanilla VEREFOO with the ability to reuse an existing valid configuration when the NSR set changes, rather than recomputing everything from scratch [3]. It takes two NSR sets as input: the *initial* set P_0 , representing the requirements already enforced by the current configuration, and the *target* set R_t , representing the new requirements that must be incorporated. Internally, React-VEREFOO classifies every NSR into deleted (R_d), added (R_a), or kept (R_k) groups (as described in Chapter 2), identifies the minimal network area affected by the added requirements, and restricts the solver to that area while leaving the rest of the configuration untouched.

Algorithm 2 React-VEREFOO

Require: Current graph G , initial NSR set P_0 , target NSR set R_t
Ensure: Updated firewall allocation satisfying all $p \in P_0 \cup R_t$, or UNSAT

- 1: Classify NSRs: R_k, R_a, R_d
- 2: Identify reconfigurable area of G based on R_a
- 3: $result \leftarrow \text{SOLVEMAXSMT}(G, initial = P_0, added = R_a)$
- 4: **if** $result = \text{SAT}$ **then**
- 5: Extract and return updated firewall allocation and filtering rules
- 6: **else**
- 7: **return** UNSAT
- 8: **end if**

React-VEREFOO issues a single MaxSMT call, just like vanilla VEREFOO, but the problem it encodes is smaller: only the nodes in the reconfigurable area appear as free decision variables. The efficiency gain therefore depends directly on how large R_a is relative to the total NSR set — the smaller the added set, the fewer nodes are flagged, and the smaller the solver problem. This relationship is captured by the percentage of requirements kept: when most requirements are kept, React-VEREFOO is significantly faster than vanilla; when most requirements change, the benefit in execution time is reduced.

At the implementation level, React-VEREFOO is invoked through the same `VerefooSerializer` constructor as the vanilla variant, with an additional boolean flag that enables the reconfiguration mode:

Listing 4.2: React-VEREFOO constructor.

```
public VerefooSerializer(NFV root, String algo, Boolean react) {
    ...
}
```

The iterative approach reuses the same reconfiguration primitive as React-VEREFOO, but rather than invoking it once with a large set of new requirements,

it invokes it repeatedly with $|R_a| = 1$ at each step, regardless of how many new NSRs must ultimately be incorporated.

4.1.3 Iterative VEREFOO

The iterative variant builds on React-VEREFOO as an internal primitive introducing NSRs one at a time. The key insight is that at every iteration the solver receives a minimal incremental problem with only one new requirement on top of a stable and already-verified configuration.

The algorithm maintains two evolving structures across iterations: *consideredProps*, the set of NSRs already enforced and passed as the *initial* set to React-VEREFOO and *currentGraph*, the network graph as updated by the most recent SAT result. Both are initialised before the loop begins: *consideredProps* with any pre-existing NSRs (denoted P_0), and *currentGraph* with the original input graph.

Algorithm 3 Iterative VEREFOO

Require: Service Graph G , pre-existing NSR set P_0 , new NSR sequence $P = (p_1, p_2, \dots, p_n)$

Ensure: Firewall allocation and filtering rules satisfying all $p \in P_0 \cup P$, or UNSAT at the failing step

```

1: Build Allocation Graph from  $G$ 
2:  $consideredProps \leftarrow P_0$ 
3:  $currentGraph \leftarrow G$ 
4: for each  $p_i \in P$  do
5:    $result \leftarrow \text{SOLVEMAXSMT}(currentGraph, initial = consideredProps, added = \{p_i\})$ 
6:   if  $result = \text{SAT}$  then
7:      $consideredProps \leftarrow consideredProps \cup \{p_i\}$ 
8:      $currentGraph \leftarrow$  updated graph from  $result$ 
9:   else
10:    return UNSAT at step  $i$ 
11:   end if
12: end for
13: return firewall allocation and filtering rules from  $currentGraph$ 

```

The total number of MaxSMT calls equals $n = |P|$, one per new NSR. Each call operates on a problem whose *added* set contains exactly one requirement — the smallest possible incremental problem. The *initial* set grows by one element at each step, shifting requirements from added to kept as the loop progresses.

4.1.4 Structural Comparison

The three variants differ along several structural dimensions, summarised in Table 4.1. The most significant difference is the number of solver invocations: both vanilla VEREFOO and React-VEREFOO issue a single MaxSMT call, while the

iterative variant issues one call per new NSR. Consequently, the scope of each call changes: vanilla encodes the full NSR set at once, React-VEREFOO processes a batch of added requirements against the existing configuration, and each iterative call encodes exactly one new requirement on top of a fixed base. Finally, vanilla VEREFOO produces a globally optimal solution, while both React-VEREFOO and the iterative variant are locally optimal — they optimise within the identified reconfigurable area but cannot guarantee global optimality across the full NSR set.

Table 4.1: Structural comparison of the three VEREFOO variants.

	Vanilla VEREFOO	React- VEREFOO	Iterative VEREFOO
Solver invocations	1	1	one per new NSR
Added NSRs per call	all at once	variable (batch)	always 1
Reuses existing config	no	yes	yes (at each step)
Optimality	globally	locally	locally (greedy)

4.2 Worked Example

To show the mechanics of the iterative approach concretely, this section traces a complete execution on a small toy example network. The scenario uses whitelisting policy mode, in which all traffic is blocked by default and only flows covered by an explicit ALLOW rule are permitted.

4.2.1 Network Topology and NSR Set

The network consists of three endpoint nodes: a *Client* (C), a *WebServer* (W), and a *Database* (D), connected in a linear chain. Traffic from C to D must therefore transit through W. Two Allocation Places — AP_1 on the C–W link and AP_2 on the W–D link — are the candidate positions where the solver may allocate firewalls.




Figure 4.1: Simple example topology. Solid circles represent endpoints; dashed rectangles are Allocation Places.

Three NSRs must be enforced, introduced one at a time in the following order:

- $p_1: C \rightarrow W$ — *reachability* (client must be able to reach the web server)
- $p_2: W \rightarrow D$ — *reachability* (web server must be able to access the database)
- $p_3: C \rightarrow D$ — *isolation* (client must not be able to reach the database directly)

The initial pre-existing set is $P_0 = \emptyset$: no firewall is allocated and no requirement is enforced at the start.

4.2.2 Iteration 1 — Introducing p_1

- $initial = \emptyset, \quad added = \{p_1\}$
- $R_k = \emptyset, \quad R_a = \{p_1\}, \quad R_d = \emptyset$

React-VEREFOO analyses the path $C \rightarrow W$ and identifies AP_1 as the relevant allocation area. Because all traffic is blocked by default and p_1 requires the path $C \rightarrow W$ to be permitted, the solver allocates a firewall FW_1 at AP_1 with the following filtering policy:

FW_1 in AP_1 : `ALLOW C → W` | `default: DROP`

After the solver returns SAT, p_1 is moved into *consideredProps* and the graph is updated to G_1 .

4.2.3 Iteration 2 — Introducing p_2

- $initial = \{p_1\}, \quad added = \{p_2\}$
- $R_k = \{p_1\}, \quad R_a = \{p_2\}, \quad R_d = \emptyset$

The reconfigurable area is now the path $W \rightarrow D$, which involves AP_2 . The firewall FW_1 at AP_1 is kept untouched, as it lies outside the identified area. The solver allocates a second firewall FW_2 at AP_2 :

FW_1 in AP_1 : `ALLOW C → W` | `default: DROP`
 FW_2 in AP_2 : `ALLOW W → D` | `default: DROP`

After the execution of this iteration, p_2 is added to *consideredProps* and the graph is updated to G_2 .

4.2.4 Iteration 3 — Introducing p_3

- $initial = \{p_1, p_2\}$, $added = \{p_3\}$
- $R_k = \{p_1, p_2\}$, $R_a = \{p_3\}$, $R_d = \emptyset$

p_3 requires that C cannot reach D. The only path from C to D passes through AP_1 . At AP_1 , FW_1 has a single ALLOW rule for traffic whose destination is W; any packet from C destined for D carries a different destination address and therefore hits the default DROP action. The isolation requirement is already satisfied by the current configuration without any modification.

React-VEREFOO confirms the existing configuration as SAT and no new firewall or rule is added. p_3 is moved into *consideredProps* and the graph remains G_2 .

4.2.5 Final Configuration and Summary

The final allocation consists of two firewalls, one at each allocation place, with minimal rule sets:

FW_1 in AP_1 : ALLOW $C \rightarrow W$ | default: DROP
 FW_2 in AP_2 : ALLOW $W \rightarrow D$ | default: DROP

Table 4.2 summarises the state of all sets at each iteration.

Table 4.2: Evolution of NSR sets and PercReqKept across the three iterations.

Iter.	Added	R_k	R_a	R_d
1	p_1	$\emptyset$	$\{p_1\}$	$\emptyset$
2	p_2	$\{p_1\}$	$\{p_2\}$	$\emptyset$
3	p_3	$\{p_1, p_2\}$	$\{p_3\}$	$\emptyset$

Three observations follow from this toy example. First, as more requirements are incorporated, the area of the network that the solver must reconsider is smaller at each iteration: the nodes satisfying all previously introduced requirements are treated as fixed, and only the area relevant to the single new requirement is left as a free decision variable. Secondly, already-enforced requirements impose no additional solving cost since their satisfaction is guaranteed by construction from the previous step and the solver does not need to reason about them again. Finally, iteration 3 required no modification at all as the isolation requirement was already implied by the whitelisting default of FW_1 , showing that the iterative approach can incorporate new NSRs with zero reconfiguration cost when the existing configuration already covers them.

These observations are specific to this toy network and cannot be generalised directly to real deployments, where larger topologies and more complex requirement interactions make the outcome of each iteration far less predictable. They

are presented here purely as a reasoning tool, to build intuition for why the iterative approach should, in principle, reduce the overall solving effort compared to a monolithic invocation. Whether this translates into a measurable gain in practice is the question addressed experimentally in Chapter [5](#).

Chapter 5

Implementation and Validation

5.1 Implementation

The iterative variant is implemented by modifying `VerefooSerializer`, the class that serves as the entry point to the VEREFOO tool: all invocations — vanilla, React, and now iterative — start here. Rather than passing the full set of new requirements to React-VEREFOO in a single call, the modified constructor wraps the React-VEREFOO invocation in a loop that introduces one NSR at a time. The vanilla constructor, also defined in `VerefooSerializer`, is left untouched, as are all downstream components — `VerefooProxy`, the MaxSMT encoding, and the `Translator`.

5.1.1 Entry Point and initialization

Since the iterative variant modifies the React-VEREFOO constructor, it shares the same parameters: `root` is the input NFV object containing the network topology and the full NSR set, `algo` selects the algorithm variant ("AP" for Atomic Predicates or "MF" for MaxFlow)¹, and `react` is the boolean flag that activates the React-VEREFOO reconfiguration.

The constructor begins by building the Allocation Graph from the input Service Graph, exactly as the other two variants do:

Listing 5.1: Allocation graph construction.

```
AllocationGraphGenerator agg = new
    AllocationGraphGenerator(root, react);
root = agg.getAllocationGraph();
```

Before the iterative loop starts, two data structures are initialized: `consideredProps` and `currentGraph`. `consideredProps` holds the requirements that have already been enforced and which will be passed as the *initial* set to React-VEREFOO at

¹In this thesis, only the Atomic Predicates algorithm is used, as React-VEREFOO supports only the AP variant.

each step, while `currentGraph` holds the current state of the network graph, which is updated after every successful solver invocation.

Listing 5.2: Initialization of the iterative state.

```
List<Property> consideredProps = new ArrayList<>();
if (initialProp != null) {
    consideredProps.addAll(initialProp);
}
Graph currentGraph = g;
```

If the input NFV object contains a pre-existing set of already enforced requirements (`initialProp`), they are loaded into `consideredProps` at the start. This supports the general case in which the network already has a valid configuration before the first new NSR is introduced, corresponding to $P_0 \neq \emptyset$ in the algorithm description of Chapter 4.

5.1.2 Core change: the Iterative Loop

The main loop iterates over `prop`, the list of new requirements (NSRs) to incorporate. At each step, `VerefooProxy` expects two NSR lists: `consideredProps` as the initial set of requirements already enforced, and `singleProp` as the new property to be added². `VerefooProxy` then computes the difference between the two lists to derive the *added*, *kept*, and *deleted* requirement sets, as described in Chapter 2. By construction, R_a (the *added* set) always contains exactly one element.

Listing 5.3: Core of the iterative loop.

```
for (Property currentProp : prop) {
    List<Property> singleProp = new ArrayList<>();
    singleProp.add(currentProp);
    singleProp.addAll(consideredProps);

    VerefooProxy test = new VerefooProxy(
        currentGraph,
        root.getHosts(),
        root.getConnections(),
        root.getConstraints(),
        consideredProps, // initial set (already enforced)
        singleProp, // new property to be enforced
        paths,
        AlgoUsed
    );
    ...
}
```

²`singleProp` is built as shown in Listing 5.3 to match the input format expected by the `VerefooProxy` interface.

After each `VerefooProxy` call, the result is examined. If all the requirements are satisfied, the current requirement is added to `consideredProps` and `currentGraph` is updated to the graph produced by the solver, so that the intermediate result is propagated to the next iteration. If the solver returns UNSAT, the loop does not break immediately, but instead the current `root` is kept as the result and the failure is recorded and surfaced to the caller after the loop completes.

5.1.3 Summary of Changes

In summary, the core change to the original React-VEREFOO is the introduction of the iterative loop: instead of passing the full set of new requirements to React-VEREFOO in a single call, the execution is broken in multiple intermediate steps, each adding exactly one requirement at a time. Between steps, the graph and the set of enforced requirements are updated, so that each React-VEREFOO invocation builds on a valid and already verified configuration. All other components — `VerefooProxy`, the MaxSMT encoding, the `Translator` — remain unchanged.

5.2 Testing

5.2.1 Qualitative Validation on CESNET

To validate the correctness of the iterative algorithm on a realistic topology, this section traces its execution step by step on the CESNET network — the Czech academic and research backbone shown in Figure 5.1. The test is qualitative: NSRs and an initial firewall configuration are defined on the CESNET real topology, and each iteration of the algorithm is applied by hand to verify that (i) each new requirement is satisfied, and (ii) no previously satisfied requirement is broken. This mirrors the structure of the Worked Example in Chapter 4, applied to a real-world topology.




Figure 5.1: CESNET backbone topology. Nodes represent Czech cities interconnected by high-speed academic links.

Topology

The topology used for this testing case is the backbone of CESNET, the Czech National Research and Education Network, shown in Figure 5.1. As modeled here it consists of 23 nodes and 28 links connecting the country’s main cities and research sites. Link capacities go from 10 Mbit/s to 10 Gbit/s, reflecting a two-tier structure: a small set of high-capacity backbone links between major cities and lower-capacity links reaching smaller, peripheral sites. Praha acts as the central node having six direct neighbors; every other node is reached through one or more intermediate hops, and some destinations are reachable through more than one path.

Initial Configuration (P_0)

This is a plausible test case constructed to exercise the iterative implementation, not an actual CESNET deployment: node names and the physical topology (links and capacities) are taken from the real network, but the reachability and isolation requirements are defined for this evaluation and do not reflect any real security policy in use on CESNET.

Ten initial NSRs are defined (P_0). Praha requires bidirectional reachability with Brno, Plzen, and Usti n.L. (p_1 – p_6). Ostrava must not be able to reach Praha, Plzen, or Usti n.L. (p_7 – p_9). Brno must not be able to reach Plzen (p_{10}). Table 5.1 lists the ten requirements.

Table 5.1: Initial NSR set P_0 (10 requirements already enforced).

ID	Source	Destination	Type
p_1	Praha	Brno	Reachability
p_2	Brno	Praha	Reachability
p_3	Praha	Plzen	Reachability
p_4	Plzen	Praha	Reachability
p_5	Praha	Usti n.L.	Reachability
p_6	Usti n.L.	Praha	Reachability
p_7	Ostrava	Praha	Isolation
p_8	Ostrava	Plzen	Isolation
p_9	Ostrava	Usti n.L.	Isolation
p_{10}	Brno	Plzen	Isolation

Firewall placement. P_0 is the initial set of ten NSRs listed in Table 5.1, and the firewall configuration described below — three firewalls — is taken as the starting point from which the iterative case unfolds. Two of them isolate Ostrava, as required by p_7 – p_9 : FW_1 closes off Ostrava’s direct routes toward Praha, Plzen, and Usti n.L. to the north, while FW_2 , on the Brno–Praha link, closes off the direct route south through Brno. The remaining one, FW_3 , closes the one route out of Ostrava that neither FW_1 nor FW_2 reaches, via Brno, Jihlava, and beyond — and, at the same time, is exactly what p_{10} needs, since it also sits on Plzen’s southern access link, so both requirements are covered at no extra cost. The three firewalls are summarized below, followed by a more detailed walkthrough of each. Every reachability requirement, instead, is already satisfied by unguarded links, so no firewall is needed to enforce them — this includes p_3/p_4 : the Mar.Lazne–Plzen link never needs a firewall, because p_{10} ’s only offending traffic (from Brno) is already stopped upstream at FW_2 and FW_3 , before it could ever reach Plzen from that direction.

- FW_1 at $AP(\text{Hradec Kralove, Letohrad})$:
 default: DROP (p_7, p_8, p_9)
- FW_2 at $AP(\text{Brno, Praha})$:
 ALLOW Praha $\rightarrow$ Brno (p_1)
 ALLOW Brno $\rightarrow$ Praha (p_2)
 default: DROP (p_7, p_8, p_9, p_{10})
- FW_3 at $AP(\text{C.Budejovice, Plzen})$:
 default: DROP (p_7, p_8, p_9, p_{10})




Figure 5.2: Subset of the CESNET topology relevant to P_0 : reachability paths (p_1 – p_6) in green, isolation paths (p_7 – p_{10}) in red, and the three firewalls allocated to enforce them.

Incremental Additions

From this point on, every step is driven by Iterative React-VEREFOO logic (Chapter 2): Step 1 classifies each new NSR as a reachability or an isolation requirement. Step 2 then differs between the two requirement types: for a reachability requirement, it checks whether some flow already reaches the destination — if so, no further action is needed; otherwise, every node currently blocking that flow becomes eligible for reconfiguration. For an isolation requirement, it checks whether the flow is already blocked — if not, every node on the path becomes a candidate for reconfiguration. Step 3 then solves a MaxSMT problem restricted to the flagged candidates eligible for reconfiguration, reusing existing firewalls and rules wherever possible.

Five new NSRs are introduced and summarized in Table 5.2.

Table 5.2: 5 additional requirements to be added iteratively.

ID	Source	Destination	Type
p_{11}	Podebrady	Usti n.L.	Isolation
p_{12}	Tabor	Pardubice	Isolation
p_{13}	Opava	Usti n.L.	Reachability
p_{14}	Olomouc	Praha	Isolation
p_{15}	Tabor	Usti n.L.	Reachability

Iteration 1 — p_{11} : Pódebrady $\nrightarrow$ Ustí n.L. (Isolation).

$$R_k = P_0, \quad R_a = \{p_{11}\}, \quad R_d = \emptyset$$

Traffic Pódebrady $\rightarrow$ Ustí n.L. would follow Pódebrady-Praha-Ustí n.L.. Nothing currently blocks it: both $AP(\text{Praha}, \text{Pódebrady})$ and $AP(\text{Praha}, \text{Ustí n.L.})$ are plain forwarders. Per Step 2, both links on the path become candidate Allocation Places for hosting the new isolation. The solver must choose one:

- Placing the firewall at $AP(\text{Praha}, \text{Ustí n.L.})$ would also cut the Praha $\leftrightarrow$ Ustí n.L. traffic already enforced by p_5/p_6 , requiring two compensating ALLOW rules in addition to the default DROP.
- Instead, placing it at $AP(\text{Praha}, \text{Pódebrady})$ intercepts no other requirement and a simple default DROP rule appears sufficient.

Biased toward the minimal number of rules, React-VEREFOO allocates the cheaper option by choosing to allocate a new firewall at $AP(\text{Praha}, \text{Pódebrady})$:

FW_4 at $AP(\text{Praha}, \text{Pódebrady})$:

default: DROP

(p_{11} , new)

No existing firewall is touched.

Result: SAT. New firewall allocated.




Figure 5.3: Iteration 1: p_{11} (Pódebrady $\nrightarrow$ Ustí n.L.) is enforced by allocating FW_4 at $AP(\text{Praha}, \text{Pódebrady})$.

Iteration 2 — p_{12} : Tabor $\nrightarrow$ Pardubice (Isolation).

$$R_k = P_0 \cup \{p_{11}\}, \quad R_a = \{p_{12}\}, \quad R_d = \emptyset$$

Traffic Tabor $\rightarrow$ Pardubice can follow two distinct routes: the direct **Tabor-Praha-Pardubice** link, but also the longer route which reaches Pardubice from its Hradec Kralove side. Nothing currently blocks either route, so Step 2 flags every link on both paths as a candidate. Since Tabor is leaf node with only one physical link, the obvious choice is to place the additional firewall on that link, namely $AP(\text{Praha}, \text{Tabor})$. Placing the firewall there closes both routes with a single default DROP rule, at half the cost of any other candidate which would had need another firewall to block the other route.

FW_5 at $AP(\text{Praha}, \text{Tabor})$:

default: DROP

(p_{12} , new)

Result: SAT. New firewall allocated.




Figure 5.4: Iteration 2: p_{12} (Tabor $\nrightarrow$ Pardubice) is enforced by allocating FW_5 at $AP(\text{Praha}, \text{Tabor})$, the only candidate closing both routes to Pardubice at once.

Iteration 3 — p_{13} : Opava $\rightarrow$ Usti n.L. (Reachability).

$$R_k = P_0 \cup \{p_{11}, p_{12}\}, \quad R_a = \{p_{13}\}, \quad R_d = \emptyset$$

Traffic from Opava toward Usti n.L. would follow **Opava-Letohrad-Hradec Kralove-Liberec-Usti n.L.**, crossing FW_1 at $AP(\text{Hradec Kralove}, \text{Letohrad})$, which

currently carries only a bare default DROP with no exceptions. The flow is blocked there, so FW_1 is the (only) candidate flagged by Step 2. This does not conflict with the standing isolation of Ostrava (p_7-p_9). React-VEREFOO reconfigures FW_1 :

FW_1 updated ($AP(\text{Hradec Kralove, Letohrad})$):
ALLOW Opava $\rightarrow$ Usti n.L. (p_{13} , new)
 default: DROP

Ostrava, Olomouc, and every other pair crossing this link still have no matching ALLOW rule, so p_7-p_9 remain fully enforced.

Result: SAT. FW_1 reconfigured.




Figure 5.5: Iteration 3: p_{13} (Opava $\rightarrow$ Usti n.L.) is enforced by reconfiguring FW_1 at $AP(\text{Hradec Kralove, Letohrad})$ with an ALLOW exception.

Iteration 4 — p_{14} : Olomouc $\nrightarrow$ Praha (Isolation).

$$R_k = P_0 \cup \{p_{11}, p_{12}, p_{13}\}, \quad R_a = \{p_{14}\}, \quad R_d = \emptyset$$

Olomouc has three routes toward Praha, and all three cross a firewall already in place. The first goes through FW_1 (via Letohrad and Hradec Kralove); the second, through FW_2 at $AP(\text{Brno, Praha})$ (via Zlín and Brno); the third — via Ostrava, Brno, Jihlava, and beyond, the same bypass identified for p_7-p_9 in P_0 — through FW_3 at $AP(\text{C.Budejovice, Plzen})$. None of the three firewalls has an ALLOW rule matching this pair. Every route is therefore already blocked, and Step 2 finds nothing to reconfigure.

Result: SAT, no change.

Iteration 5 — p_{15} : Tabor $\rightarrow$ Usti n.L. (Reachability).

$$R_k = P_0 \cup \{p_{11}, p_{12}, p_{13}, p_{14}\}, \quad R_a = \{p_{15}\}, \quad R_d = \emptyset$$

Traffic Tabor $\rightarrow$ Usti n.L. would follow Tabor-Praha-Usti n.L., crossing FW_5 at $AP(\text{Praha}, \text{Tabor})$, which since iteration 2 carries only a default DROP with no exceptions. The flow is blocked there, so FW_5 is the (only) candidate flagged by Step 2 needed to be reconfigured. This does not conflict with the standing isolation of Pardubice (p_{12}). React-VEREFOO reconfigures FW_5 :

FW_5 updated ($AP(\text{Praha}, \text{Tabor})$):

ALLOW Tabor $\rightarrow$ Usti n.L.

(p_{15} , new)

default: DROP

Pardubice and every other destination reached through this link still have no matching ALLOW rule, so p_{12} remains fully enforced.

Result: SAT. FW_5 reconfigured.




Figure 5.6: Iterations 4 and 5: p_{14} (Olomouc $\nrightarrow$ Praha) is already blocked by FW_1 , FW_2 , and FW_3 on every route, while p_{15} (Tabor $\rightarrow$ Usti n.L.) is blocked by FW_5 and requires adding an ALLOW exception to it.

Final Configuration Verification

All five new NSRs were successfully incorporated. The network ends with five firewalls: the three from P_0 (FW_1 reconfigured in iteration 3) plus FW_4 and FW_5 , allocated at $AP(\text{Praha}, \text{Podebrady})$ and $AP(\text{Praha}, \text{Tabor})$ respectively, with FW_5

later reconfigured in iteration 5. Table 5.3 verifies each of the 15 enforced requirements against the final configuration: at the end of every iteration, all requirements in `consideredProps` are simultaneously satisfied.

Table 5.3: Verification of all 15 NSRs against the final firewall configuration.

ID	Type	Enforcement mechanism	FW
p_1	Reach.	ALLOW Praha $\rightarrow$ Brno	FW_2
p_2	Reach.	ALLOW Brno $\rightarrow$ Praha	FW_2
p_3	Reach.	Unfiltered link (no firewall needed)	—
p_4	Reach.	Unfiltered link (no firewall needed)	—
p_5	Reach.	Unfiltered link (no firewall needed)	—
p_6	Reach.	Unfiltered link (no firewall needed)	—
p_7	Isol.	DROP by default on every route out of Ostrava	FW_1, FW_2, FW_3
p_8	Isol.	DROP by default on every route out of Ostrava	FW_1, FW_2, FW_3
p_9	Isol.	DROP by default on every route out of Ostrava	FW_1, FW_2, FW_3
p_{10}	Isol.	DROP by default on every route from Brno to Plzen	FW_2, FW_3
p_{11}	Isol.	(iteration 1) DROP by default at FW_4	FW_4
p_{12}	Isol.	(iteration 2) DROP by default at FW_5	FW_5
p_{13}	Reach.	(iteration 3) ALLOW Opava $\rightarrow$ Usti n.L.	FW_1
p_{14}	Isol.	(iteration 4) DROP by default on every route from Olomouc	FW_1, FW_2, FW_3
p_{15}	Reach.	(iteration 5) ALLOW Tabor $\rightarrow$ Usti n.L.	FW_5

Discussion

The clearest takeaway, already visible in P_0 , is that the *type* of NSR — reachability or isolation — is what really drives the outcome of the algorithm: the two types make it iterate in completely different ways. A reachability requirement can only find its path already unprotected (as in iteration 5) or already being blocked by an existing firewall (resulting in reconfiguration as in iteration 3), but will never justify allocating a new firewall. An isolation requirement, instead, can force a genuinely new allocation, to start blocking traffic where nothing did before (iterations 1 and 2); but it can also find every path already closed by firewalls allocated for earlier requirements, needing no change at all (iteration 4). This asymmetry — reachability at most reconfiguring an existing firewall, isolation possibly allocating a new one — is itself a consequence of the *whitelisting* mode chosen for this testing case (Chapter 2).

Iterations 1 and 2 both allocate a new firewall for an isolation requirement, but each selects its Allocation Place for a different reason. In iteration 1 the two direct candidates have different cost — one is free, the other would break an existing reachability requirement — so the cheaper option is the clear choice. Iteration 2 looks similar at first, since Tabor and Pardubice appear nowhere else in the NSR set, but the isolation itself is reachable by two distinct routes, not only one: closing

only the direct Tabor–Pardubice link leaves the longer route through Usti n.L. open. $AP(\text{Praha}, \text{Tabor})$ is the only candidate that sits on both routes at once — since Tabor is a leaf node linked to Praha — so it closes the isolation with a single firewall, at half the cost of any alternative. Allocating a firewall on link $AP(\text{Praha}, \text{Tabor})$ forces a reconfiguration at iteration 5 to allow traffic from Tabor to Usti n.L., as required by p_{15} .

The final result counts five firewalls enforcing all fifteen requested NSRs. The step-by-step verification carried out at each iteration confirms that the requirements introduced during the incremental phase were correctly enforced, without disrupting any requirement already in place.

5.2.2 A Closer Look: Tie-Breaking and Flow Analysis

When React-VEREFOO must allocate a new firewall to satisfy an isolation requirement, it is not unusual for two or more Allocation Places to turn out equally costly with respect to the requirement being processed: assuming each candidate, would need the same number of new rules, there is no reason to prefer one over the other at that stage. But it may happen that this is not neutrally propagated into the following iterations. If a later NSR happens to require reachability across that same link, Step 2 of the iterative process will find it blocked by the firewall just placed, and will need to reconfigure it with a new ALLOW exception. However, in an alternative scenario where the equally-cheap candidate would have been chosen, the later requirement might have found its path still open at later iteration, and been satisfied at zero cost. Both outcomes remain correct regardless of which candidate is chosen: the difference only relies on the number of configurations rules needed, not on their validity.

This observation - taken as an example of several equally plausible scenarios that could arise elsewhere in the process - points to a natural extension of the approach, that falls outside the scope of this thesis: a *traffic-aware* strategy may be considered for future work, applied before listing the NSRs and presenting them to the iterative version of React-VEREFOO, aiming to classify NSRs by their traffic flows and order them across iterations so as to minimize the number of future reconfigurations, or to break such ties by anticipating which candidate is less likely to be needed again later.

5.2.3 Discussion and Limitations


A monolithic benchmark. To put the iterative result in context, vanilla VEREFOO was also run once, monolithically, on the very same CESNET topology and the complete set of fifteen NSRs — the ten of P_0 together with the five introduced across the incremental phase — handed to the solver all at once instead of one at a time. The topology was modeled with the same conventions used throughout this validation: every router is a fixed forwarder, and each of the 28 links carries its own dedicated Allocation Place, exactly as in the iterative trace. This ensures the comparison only refers to *when* the solver sees the requirements, whether all at once or one at a time.

The monolithic run is SAT, and it satisfies all fifteen requirements with five firewalls, exactly as many as the iterative trace. Table 5.4 summarizes the result.

Table 5.4: Firewalls allocated by a single, monolithic run of vanilla VEREFOO over all 15 NSRs.

Firewall	Allocation Place	Enforcement mechanism	NSRs covered
FW_1^M	$AP(\text{Praha, Brno})$	DROP by default; ALLOW $\text{Praha} \leftrightarrow \text{Brno}$	p_1, p_2
FW_2^M	$AP(\text{Praha, P\odebrady})$	DROP by default	p_{11}
FW_3^M	$AP(\text{Praha, Pardubice})$	DROP by default	$p_{12}, p_7\text{--}p_9,$ p_{14}
FW_4^M	$AP(\text{Usti n.L., Liberec})$	DROP by default; ALLOW $\text{Opava} \rightarrow \text{Usti n.L.}$	p_{13}
FW_5^M	$AP(\text{C.Budejovice, Plzen})$	DROP by default	p_8

Three of the five firewalls match exactly between the two configurations, each pair sitting on the same Allocation Place and playing the same role: FW_1^M acts exactly as FW_2 does in the iterative version, at $AP(\text{Praha, Brno})$; FW_2^M as FW_4 , at $AP(\text{Praha, P\odebrady})$; and FW_5^M as FW_3 , at $AP(\text{C.Budejovice, Plzen})$. The remaining two placements differ, and the difference traces back to a single decision: where to enforce Ostrava and Olomouc’s northern isolation ($p_7\text{--}p_9, p_{14}$). The iterative solution commits to this by placing FW_1 in $AP(\text{Hradec Kralove, Letohrad})$ already in P_0 , long before p_{12} or p_{13} exist. The monolithic run, seeing the full set of fifteen NSRs at once, instead places a firewall at $AP(\text{Usti n.L., Liberec})$ — a different point on the very same path (Opava–Letohrad–Hradec Kralove–Liberec–Usti n.L.). This is not an equivalent choice: $AP(\text{Usti n.L., Liberec})$ also happens to sit on Tabor’s alternate route to Pardubice (Tabor–Praha–Usti n.L.–Liberec–Hradec Kralove–Pardubice), so placing the isolation firewall there closes that route too, at no extra cost. This is why the monolithic solution only needs one more firewall, at $AP(\text{Praha, Pardubice})$, to close p_{12} ’s direct route, and never needs a firewall dedicated to Tabor at all. The iterative trace, having already committed FW_1 to Hradec Kralove–Letohrad in P_0 , cannot retroactively move it once p_{12} introduces the alternate-route requirement, and must instead allocate a fifth firewall dedicated to $AP(\text{Praha, Tabor})$ to close both of Tabor’s routes at once. While the two solutions are equivalent in firewall count, the iterative solution carries one extra ALLOW rule, needed in iteration 5 to let Tabor $\rightarrow$ Usti n.L. traffic through, after it was blocked by allocating FW_5 in iteration 2.



avoid this by design: once a firewall is placed, it is left untouched unless the new requirement genuinely cannot be satisfied without altering it.

The other expected advantage is scalability, though this validation is too small to show it directly: with only fifteen NSRs on this topology, both approaches solve quickly, and the small scale of this qualitative test does not stress the difference. The gain is expected to matter for larger changes — bigger topologies and larger, more frequent NSR sets — where re-solving the whole configuration from scratch grows increasingly expensive compared to solving one requirement at a time against a mostly-fixed configuration. Confirming this with a timing benchmark at scale is left as future work.

Chapter 6

Conclusion and Future Work

6.1 Conclusions

Manual firewall configuration does not scale with the size and complexity of modern networks and offers no guarantee of correctness, which motivated VEREFOO as an automated solution for configuring firewalls in complex topologies. But because modern networks are dynamic and change frequently, VEREFOO’s monolithic strategy scales poorly; this motivates React-VEREFOO’s incremental approach of reconfiguring only the portion of the network affected by the NSRs that actually changed.

This thesis pushed that same incremental principle to its extreme: rather than presenting React-VEREFOO with a batch of new requirements, the proposed Iterative React-VEREFOO introduces exactly one NSR at a time, always keeping each individual subproblem as small as it can possibly be, at the cost of turning one solver call into a chain of many smaller ones.

The approach was validated qualitatively on the CESNET academic topology. The resulting configuration was also compared against a single monolithic run of vanilla VEREFOO over the same requirements: the two solutions allocated the same number of firewalls, differing only in the placement of two of them, and the iterative configuration required one additional rule to make up for it — a minor gap, given the full set of NSRs to enforce. This shows the iterative process reaches a configuration correct and nearly as compact as the monolithic one, even though it never has visibility beyond the single requirement being processed at each step — though, as discussed next, this remains to be confirmed at a larger scale.

6.2 Future Work

The most important next step is a quantitative evaluation to measure the execution time across networks and NSR sets of varying scale, and evaluate if the iterative approach is faster than the monolithic baseline, complemented by a systematic comparison of final firewall and rule counts to quantify any trade-off in configuration quality.

Additionally, two directions could improve the approach itself. Section [5.2.2](#) showed that the order in which NSRs are introduced can affect the cost of later requirements when the solver faces a tie; a traffic-aware ordering could reduce future reconfigurations. Separately, this thesis only handles the addition of new NSRs — extending it to also support removals, and arbitrary mixes of the two, would cover the full range of changes a live network undergoes.

Bibliography

- [1] D. Brighenti, G. Marchetto, R. Sisto, and F. Valenza, “Automation for Network Security Configuration: State of the Art and Research Trends,” *ACM Computing Surveys*, vol. 56, no. 3, 2023.
- [2] D. Brighenti, G. Marchetto, R. Sisto, F. Valenza, and J. Yusupov, “Automated Firewall Configuration in Virtual Networks,” *IEEE Transactions on Dependable and Secure Computing*, vol. 20, pp. 1559–1576, March-April 2023.
- [3] F. Pizzato, D. Brighenti, R. Sisto, and F. Valenza, “Automatic and Optimized Firewall Reconfiguration,” in *Proc. IEEE/IFIP Network Operations and Management Symposium (NOMS 2024)*, 2024.
- [4] D. Brighenti, S. Bussa, R. Sisto, and F. Valenza, “Atomizing Firewall Policies for Anomaly Analysis and Resolution,” *IEEE Transactions on Dependable and Secure Computing*, vol. 22, no. 3, 2025.
- [5] D. Brighenti, S. Bussa, R. Sisto, and F. Valenza, “A Two-Fold Traffic Flow Model for Network Security Management,” *IEEE Transactions on Network and Service Management*, vol. 21, no. 4, 2024.